GestaltView Agent Trainer: Comprehensive User Onboarding & Implementation Manual

1. The Strategic Shift: From Generic Chatbots to Specialized Operators

In the current enterprise landscape, the primary friction point for AI implementation is the "plausible but ungrounded" response. Generic Large Language Models (LLMs) are eloquent but operate in a vacuum, lacking the proprietary context required for high-stakes operations. The GestaltView Agent Trainer (GAT) **enforces** a fundamental shift in philosophy, moving from a standard "chatbot" to a **specialized operator**. This is achieved through a "productized floor"—a hardened, white-label AI scaffold designed to ground assistants in an organization's unique "world." By prioritizing intentional training over mindless data ingestion, the system transforms a volatile LLM into a reliable, inspectable business asset.

The "Enemy" vs. The Solution

Generic AI (The Problem)	GestaltView Agent Trainer (The Solution)
"Raw PDF Dumping": Shoveling unorganized files into a vector store.	**"PreciseContextualizedIntelligence":** Structured, multi-lane data curation.
Hallucinations: Plausible-sounding but factually incorrect "bullshit."	**Absolute Source Fidelity:** Reasoning based on buyer-owned artifacts.
"PR Drone" Voice: Generic, unaligned corporate drone behavior.	**Brand-Faithful Alignment:** Customizable vocabulary and linguistic profiles (PLK).

Experimental/Volatile: Lacks operational rigor and IP protection.	**Production-Ready Asset:** Self-contained, inspectable environment with an Import Boundary.

Core Mission Statement: "The GestaltView Agent Trainer is a standalone, white-label AI assistant scaffold designed to allow buyers to train and deploy assistants grounded in their own proprietary corpus, vocabulary, and operating context."

This shift in philosophy necessitates a structured approach to the system's core interfaces, ensuring that both technical and business stakeholders can contribute to the assistant's "world-awareness."

----

2. System Architecture & The "Jack-In" Methods

Deploying professional-grade AI requires owning the underlying infrastructure to ensure data sovereignty and intellectual property protection. The GAT is built on "Core Cybernetics"—a stack that separates configuration from backend logic, allowing operators to inspect specific layers without exposing the internal runtime logic of the parent framework.

The Technical Stack

- **The Data Vault:** Powered by **Supabase and pgvector**, this layer stores chunked proprietary data. It utilizes Row-Level Security (RLS) to ensure that users only access data linked to their specific workspace or ID.
- **The Neural Prism (LLM Router):** An agnostic provider chain that routes prompts to specific models like **Groq (Llama-3), OpenAI (GPT-4), or Gemini**. The router coordinates context assembly and enforces persona boundaries.
- **The Jack-In Interface:** A **React / Next.js** client-facing UI serving as the management dashboard and validation console.

Technical Note on Retrieval: To achieve maximum accuracy, the GAT utilizes a **Hybrid Retrieval** methodology. This combines **Vector Similarity Search (cosine similarity)** for semantic relevance with **Full-Text Search (keyword-based ranking)** to ensure the assistant catches specific technical terminology often missed by pure vector models.

Choose Your Path

The system supports two entry points for onboarding, both of which feed into a shared task graph to ensure no loss of state during the transition between technical setup and visual configuration.

Path A: GV Concierge CLI (gv.sh)	Path B: Setup Wizard
Persona: Technical Operator / "Terminal Junkie"	**Persona:** Business Operator / Visual Architect
Capabilities: Run 'doctor' checks on environment posture, execute manifest-driven JSON imports, and manage status.	**Capabilities:** Browser-based identity configuration, visual selection of Domain Presets, and theme selection.
Unique Feature: Generates "handoff bundles"—machine-readable snapshots of the session state for troubleshooting.	**Unique Feature:** Click-to-build `.env.local` file generation and visual "identity drafting."

----

3. Phase I: Constructing the Foundation & Commercial Identity

The **Import Boundary** is the heart of the GAT architecture, establishing a legal and technical isolation layer that ensures proprietary data remains sequestered from the core framework. Phase I focuses on environment-driven verticalization, where the functional "lens" of the assistant is defined through configuration variables.

Infrastructure Provisioning

The foundation is built through four essential steps:

1. **Database Provisioning:** Initializing the Supabase project and enabling the pgvector extension for high-dimensional vector embeddings.
2. **Environment Configuration:** Populating the .env.local file to "wire" the system to its data and power sources.
3. **Verification:** Executing the npm run verify-setup script to validate that all mandatory SQL migrations and task graphs are present.
4. **Identity Onboarding:** Naming the workspace, selecting a Domain Preset, and establishing visual branding tokens.

Critical Environment Variables

Variable	Purpose
`KIT_NAME`	User-defined string that establishes brand recognition and product identity.
`NEXT_PUBLIC_SUPABASE_URL`	The API endpoint for your database project.
`SUPABASE_SERVICE_ROLE_KEY`	Elevated key required for server-side operations and corpus ingestion.
`KIT_TIER`	Enforces operational boundaries, seat counts, and fragment limits.
`EMBEDDING_PROVIDER`	The service used for vectorization (e.g., OpenAI, Gemini).
`KIT_DOMAIN`	The primary trigger for `domainPrompts` logic; sets the cognitive "lens."
`KIT_PRIMARY_COLOR`	Hex code that synchronizes the UI with the client's corporate identity.

Domain Presets & Aesthetic Tiers

The GAT utilizes **Domain Presets** to define the "Starter Prompt Focus." Operators must select the appropriate lens:

- **Consulting:** Focuses on **Expert-system framing**. Suggested uploads: Service offers, playbooks, and case studies.

- **ADHD:** Focuses on **Reducing ambiguity and task scaffolding**. Suggested uploads: Task systems and checklists.
- **Creative:** Focuses on **Preserving original voice while refining ideas**. Suggested uploads: Brand voice guides and creative briefs.

Visual identity is finalized via **Theme Presets**:

- **Lagoon Glass:** Aqua gradients and a premium "glass studio" feel for creative sectors.
- **Signal Noir:** A high-contrast "command center" style for technical verticals.
- **Atlas Neutral:** A restrained, corporate look prioritizing accessibility and professional decorum.

Commercial Tier Selection

- **Solo Track:** Focused on personal productivity and rapid publication; single-seat starter.
- **Business Track:** Focused on team collaboration and reusable "Operator Packs"; shared workspaces.
- **Enterprise Track:** Focused on governance, policy enforcement, and multi-workspace control.

----

4. Phase II: Training via the "Four-Lane" Corpus Model

AI grounding is primarily a data management problem. The GAT **mandates** the "Curation over Dumping" principle: 10–25 high-signal files per lane are significantly more effective than massive repository dumps.

The Four Data Highways

Lane	Weight	Target Material	Assistant Outcome
Knowledge	30%	SOPs, FAQs, reference manuals.	A factually grounded, reliable operator.
Code	25%	READMEs, architecture notes, API specs.	A repo-aware, technical collaborator.

Product	25%	Specs, roadmaps, release notes.	A product-aware partner for strategy.
Context	20%	Voice notes, values, positioning.	An aligned, brand-faithful assistant.

The BYOK Ingestion Pipeline

1. **Raw Input:** Selection of high-signal documents for each lane.
2. **The Slicing:** Documents are sliced into **1200-character fragments** with a **200-character overlap** to prevent context loss.
3. **The Embedding:** Fragments are converted into high-dimensional vectors.
   - *Technical Note:* The system defaults to **768 dimensions for Gemini** or **1536 dimensions for OpenAI**.
4. **Storage:** Data is persisted in Supabase pgvector for hybrid similarity and keyword search.

----

5. The Personal Language Kit (PLK) & Operational Memory

To eradicate the "generic corporate drone voice," the GAT utilizes "Neural Implants" known as the Personal Language Kit. This ensures the assistant sounds native to the buyer’s organization.

Vocabulary Profile Controls

- **Preferred Language:** Direct, warm, and disciplined terminology.
- **Avoid:** Targets "forbidden filler," generic AI jargon, and inflated certainty.
- **Audience:** Tailors responses for specific groups (e.g., builders vs. knowledge-heavy teams).
- **Output Bias:** Dictates response structures (e.g., "clear framing" or "concrete next steps").

Memory Management (The Persistence Protocol)

Memory is categorized into reviewable windows to prevent "wording drift."

The Keep (Secure Storage)	The Burn (Data Destruction)
User Memory: Stable working styles and preferences.	**Secrets & Credentials:** API keys are never stored in working memory.
Shared Memory: Pinned continuity for team collaboration.	**Regulated Claims:** High-risk legal/medical data without policy review.
Pinned Memory: Durable truths and milestones that survive session noise.	**Noisy Trivia:** Useless session chatter that clogs the context window.

----

6. Phase III: Validation & The "Go-Live" Readiness Score

Deployment in the GAT is an **earned state**, not a progress bar. The "Readiness Score" serves as a mandatory quality gate for high-stakes business operations.

The 78% Threshold

An agent is only considered production-ready when it clears three strict thresholds:

- **Overall Score:** ≥ 78% (Weighted average across the four lanes).
- **Setup Completion:** ≥ 70% (Progress through the onboarding task graph).
- **Evaluation Pass Rate:** ≥ 70% (Pass rates from internal benchmarks).

Operators utilize the goLiveVerdict to manage deployment:

- **not_ready**: Critical infrastructure or data gaps present.
- **close**: Thresholds nearly met; requires final lane refinement.
- **ready**: All gates cleared; "Publish" button unlocked.

The Interrogation Room

The **Validation Console** uses high-stakes benchmark prompts to stress-test the assistant:

1. *"Summarize the current architecture and call out one likely fragility."*

2. *"Explain the product in the operator's own voice for a new customer."*
3. *"List the missing documents that would improve answer quality most."*

The **Context Stack Inspector** provides transparency, showing exactly which retrieval layers were used and providing a "Retrieval Confidence" percentage.

Activation Milestones Checklist

1. [ ] **Identity Created:** Workspace named and identity defined.
2. [ ] **LLM Provider Chain Connected:** OpenAI, Groq, or Gemini linked.
3. [ ] **Corpus Imported:** Material assigned across multiple lanes.
4. [ ] **Evaluation Run:** Benchmarks executed and pass rate recorded.
5. [ ] **Publish Approval Recorded:** Final thresholds met and approval logged.

----

7. Post-Deployment: Iterative Maintenance & Professional Services

Once the ready verdict is issued, the operator's role shifts from setup to refinement to maintain a "grounded state."

The Operator’s Rhythm

- **Trace Weak Answers:** Use **Weak Zone Tracing** to identify where retrieval is thin. If the system fails a code question, the fix is in the Code Lane, not the prompt.
- **Metadata Accuracy:** Ensure every source has a clear title and URI to provide clickable citations.
- **Curate, Don’t Dump:** One clean SOP is worth more than fifty outdated chat transcripts.

The Operator's Arsenal (Consulting Upsells)

Service Delivery Menu	Description	Commercial Value
Custom Exhibit Buildout	Creation of domain-specific UI surfaces and vocabulary profiles.	**$2,500**

Knowledge Curation	Rigorous corpus audit and fragmentation strategy for optimization.	**$150/hr**
Full Platform Deployment	End-to-end technical installation, QA, and launch support.	**$5,000**
Enterprise Engagement	High-level governance, policy planning, and scaling requirements.	**$5,000+**
The Productized Floor	The complete commercial scaffold for infinite white-label assistants.	**$2,400 (One-Time)**

----

8. Appendix: Operational Glossary & FAQ

Glossary of Key Terms

- **BYOK (Bring Your Own Knowledge):** The principle of grounding an AI assistant in proprietary data.
- **Import Boundary:** A critical legal and technical isolation layer separating buyer data from the framework engine.
- **pgvector:** A PostgreSQL extension used to store and query high-dimensional vector embeddings.
- **RLS (Row-Level Security):** Supabase policies ensuring that users only access data linked to their specific workspace.
- **Skill Fragment:** A discrete unit of behavior or instruction mapped to a specific domain (e.g., "Code Guide").
- **Domain Presets:** Specialized cognitive "lenses" that provide expert-system framing for specific industries.

Quick-Start Quiz

1. What is the primary purpose of the GestaltView Agent Trainer? It is a standalone, white-label AI scaffold designed to allow buyers to train and deploy assistants grounded in their own proprietary corpus, vocabulary, and context.

2. Explain the significance of the "Four-Lane Model" in corpus curation. It categorizes data into Knowledge, Code, Product, and Context lanes to ensure a balanced understanding of facts, architecture, roadmaps, and brand voice.

3. How does the "Go-live readiness" score determine if a workspace is ready? It requires a weighted average of ≥ 78%, setup completion ≥ 70%, and an evaluation pass rate ≥ 70%.

4. Describe the difference between the "Operator CLI" and the "Setup Wizard." The CLI (gv.sh) provides terminal-level control for technical operators, while the Setup Wizard is a browser-based tool for visual identity and environment drafting.

5. What are the three classifications of the Operational Memory model? The model includes **User Memory** (preferences), **Shared Memory** (team continuity), and **Pinned Memory** (durable truths), while excluding "Unsafe to Store" data like secrets.

Final Statement: By moving from volatile, generic LLMs to reliable, inspectable business assets, the GestaltView Agent Trainer ensures your assistant isn't just intelligent—it's an expert in your specific world.# GestaltView Agent Trainer: Comprehensive User Onboarding & Implementation Manual